



Sistemas de Operação

(Ano letivo de 2026-2027)

Guiões das aulas práticas

script #04

Processes, shared memory, and semaphores

Summary


- Understanding and dealing with concurrency using shared memory.
 - Using semaphores to control access to a shared data structure by different processes.
-

Previous note

In the code provided, some system calls are not called directly. Instead, they are called inside a macro that avoids defensive programming. The macros deal internally with error situations, either aborting execution or throwing exceptions, thus releasing the programmer of doing so. These macros will be available during the practical exams.

Exercises

Exercise 1 Implementing a bounded-buffer application, using processes, shared FIFO and semaphores.



Directory `bounded_buffer` provides an example of a simple producer-consumer application, where interaction is accomplished through a buffer with bounded capacity. The application relies on a FIFO to store the items of information generated by the producers, that can be afterwards retrieved by the consumers. Each item of information is composed of 3 integer values, one used to store the id of the producer and the other two general purpose. Directory `bounded-buffer` contains the support source code for this exercise.

(a) **Understanding the fifo data type definition**

File `fifo.h` defines a `FIFO` data type and the signature of some manipulating functions.

- Analyse it and try to understand it.

(b) **A first implementation of the fifo**

File `fifo-unsafe.cpp` implements a first version of the fifo manipulating functions, that does not take into consideration thread/concurrency safeness.

- Analyse it and try to understand the implementation of the different functions.

(c) **Understanding how the concurrency is launched**

File `main.cpp` implements the main program, which launches child processes to execute the producer and consumer life cycle procedures. Analyse it.

- Try to understand how shared memory is created and used.

- Try to understand how process creation is done.
- What is the purpose of the `exit(EXIT_SUCCESS)` after the calls to the producer and consumer procedures?
- Try to understand how the main code waits for the child processes to finish.

(d) **See race conditions showing up**

As said before, every item is composed of 3 values. When they are generated (by a producer), fields `v1` and `v2` are made equal and, in addition, they contains the `id` in their value (`v1 = v2 = id * 1000000 + i`). So, if everything goes alright, an item retrieved by a consumer must meet these restrictions. In the code provided, a consumer only prints (in red color) an item if it does meet the restrictions, meaning there was a **race condition**. The option `-V` makes all messages to appear.

- Generate the *unsafe* version of the program (`make bb-unsafe`) and execute it (`./bb-unsafe`). The execution doesn't end. Why? You can press CONTROL-C to abort the execution.
- Analyse the results of different executions. It may take several runs for the race condition effect to show up. It may even be necessary to change the number of items produced per producer. You can execute `./bb-unsafe -h` to see the command line options.
- Point out an execution scenario that might result in a race condition.
- The `Fifo` data structure should be 1216 bytes long. Command `ipcs` shows which system V resources are in use. Execute it and identify the shared memory used by the program.
- Command `ipcrm -m <shmid>` can be used to release a block of shared memory. Used it to remove the shared memory used by the program. Re-execute the `ipcs` command to see the result.
- Look again at the code of the *unsafe* version, `fifo-unsafe.cpp`, and try to understand why race conditions can appear. What should be done to solve the problem?

(e) **Understanding the safe implementation of the fifo**

Generate the *safe* version of the program (`make bb-safe`), execute it (`./bb-safe`) and analyse the results. Race conditions should no longer appear.

- Look at the code of the *safe* version, `fifo-sem-safe`, analyse it and try to understand how semaphores are used to implement the safe version of the fifo.
- What is the purpose of the `ACCESS` semaphore?
- What is the purpose of the `SLOTS` and `ITEMS` semaphores?

(f) **Training exercise 1**

In the *safe* version, the program still does not terminate. A simple form, although not entirely correct, of doing it is to kill the consumers after all the producers have terminated, and some time has passed. Alter the `main.cpp` program to accomplish it. Note that if the program get to the end, the allocated resources (shared memory and array of semaphores) are released.

(g) **Training exercise 2**

Imagine a form of clean termination and implement it. One possibility is the insertion, by the main process, after all producers' termination, of dummy items, understood by consumers as exit notifications.

Exercise 2 Implementing an up-down counter application, using a shared integer variable.

The idea is to implement a concurrent program, involving a parent process and a child process, that, in collaboration, first increment and then decrement a counter in shared memory. The conjugate behaviour should be the printing in the terminal of values from 1 to N, followed by values from N-1 to 1.

(a) The main program (parent process) should:

- create an integer variable in shared memory and start it at 1;
- launch a child process, whose functionality is given below;
- wait until the child process terminates;
- decrement the value in the shared variable until it reaches 1, printing its value at every iteration;
- release the shared memory and terminate;

(b) The child process should:

- ask the user for a value between 10 and 20, validating the value inserted;
- increment the value in the shared variable until it reaches that value, printing its value at every iteration;
- terminate.

Exercise 3 Implementing a decrementer application, using a shared integer variable.

The idea is to implement a concurrent program, involving two child processes, that, in collaboration, decrement a counter in shared memory. The conjugate behaviour should be that the shared variable is alternately decremented by the two child processes.

(a) The main program (parent process) should:


- ask the user for a value between 10 and 20, validating the value inserted;
- create an integer variable in shared memory and start it with the value read;
- launch two child processes, whose functionality is given below;
- create the semaphores required to synchronize the activity of the two child processes;
- initialize the semaphores with the appropriate values;
- wait until both child processes terminate;
- release the shared memory and the semaphores and terminate.

(b) Each child process should:

- wait until it is its turn to decrement;
- terminate if the value in the shared variable is 1.
- decrement it, if not;
- print the value saying who made the decrement (PID);
- terminate if value is 1; otherwise repeat from top.

Exercise 4 Designing and implementing a simple client-server application

The figure below represents a simplified representation of a client-server concurrent system based on shared memory. The supporting (shared) data structure consists of a pool of N buffers of communication, individually identified by a number (between 0 and $N-1$), and two fifos, one of ids of buffers available and one of ids of buffers with pending orders. The same buffer is used for a client to place a request and the server to place the response to that request.



On the client side, interaction with the server takes place according to the following pseudo-code:

```
id = getFreeBuffer();           /* take a buffer id out of fifo of free buffers */
putRequestData(data, id);       /* put request data on buffer */
submitRequest(id);              /* add buffer id to fifo of pending requests */
waitForResponse(id);            /* wait (blocked) until the response is available */
resp = getResponseData(id);     /* take response out of buffer */
releaseBuffer(id);              /* buffer is free, so add its id to fifo of free buffers */
```

On the server side, the interaction is described by the pseudo-code:

```
id = getPendingRequest();       /* take a buffer id out of fifo of pending requests */
req = getRequestData(id);       /* take the request */
resp = produceResponse(req);    /* produce a response */
putResponseData(resp, id);      /* put response data on buffer */
notifyClient(id);               /* so client is waked up */
```

This is a double producer-consumer system, requiring three types of synchronization points:

- the server must block while the fifo of pending requests is empty;
- a client must block while the fifo of free buffers is empty;
- a client must block while the response to its request is not available in the buffer.

Note that in the last case there is a synchronization point per buffer. Note also that, as long as the fifos' capacities are at least the pool capacity, there is no need for a fifo full synchronization point.

Finally, consider that the purpose of the server is to process a sentence (string) to compute some statistics, specifically the number of characters, the number of digits and the number of letters.

- Using the *safe* implementation of the fifo, used in a previous exercise, as a guideline, design and implement a safe solution to the data structure and its manipulation functions.
- Implement the server process, assuming there is only one.
- Implement the client process.
- Does your solution work if there are more than one server?